# BiasAperture — Demographic Fairness Audit Dossier

Diagnostic Demographic Bias Evaluation & Regulatory Traceability Report

AUDITED MODEL
**FairFace ResNet-34 Multi-Task Classifier**

EVALUATION DATASET
**FairFace Benchmark Dataset (N=10954)**

PROTECTED AXIS
**Race**

AUDIT DATE
**2026-10-19 04:36:51 UTC**

**Diagnostic Scope & Non-Intervention Guarantee:**

BiasAperture is engineered exclusively as an external diagnostic auditing, statistical significance testing, and regulatory traceability platform. It does **not** perform model retraining, in-processing debiasing, or synthetic image generation. In adherence to statistical integrity safeguards (NFR-003), demographic subgroups with fewer than 30 samples (*n < 30*) have point estimates suppressed to prevent unreliable compliance claims.

| 7 / 7 | Overall (Table 1) | DIR = 0.713 | 0 |
|---|---|---|---|
| SUBGROUPS WITH SIGNIFICANT DISPARITY (P < 0.05) | WORST DISPARITY COHORT (28 METRIC FAILURES) | WORST HEADLINE METRIC (OVERALL) | SUBGROUPS FLAGGED N < 30 (NFR-003) |

EXECUTIVE TRANSLATION • NON-TECHNICAL REVIEW

FAIL

## Demographic Disparity Flagged: Remediation & Review Required

The audit of FairFace ResNet-34 Multi-Task Classifier identified statistically verified demographic disparities affecting 7 of the 7 evaluated Race cohorts (N=10954 subjects, statistical significance $p < 0.05$). These performance gaps exceed random statistical fluctuation, indicating systematic differences in model predictions.

### 1 Disparity Impact

The highest demographic divergence occurred for Overall (Table 1), where the Disparate Impact Ratio (DIR) was 0.713. In plain terms, subjects in this group received positive predictions at only 71.3% the rate of the benchmark group—below the standard 80% legal threshold (a 28.7% shortfall). (Observed at the global cross-group level).

### 2 Underlying Drivers

Surrogate proxy attribution was not triggered because demographic disparities remained within statistical tolerances, or subject records were not supplied for secondary modeling.

### 3 Sample Reliability

All evaluated demographic cohorts met or exceeded the minimum sample size threshold ($n \geq 30$ subjects per NFR-003), ensuring statistical reliability across all reported metrics.

## 1. Headline Fairness Metrics (Core Four — Cross-Group Overview)

Summary metrics evaluated across all eligible demographic strata (n ≥ 30 per NFR-003). Accompanied by 95% stratified BCa bootstrap confidence intervals and Pearson's $\chi^2$ test of independence. Plotted on a calibrated [0.0, 1.0] domain with null-point reference and threshold lines.

▼ 📖 **How to Read These Metrics (In Plain English)**

### Disparate Impact Ratio (DIR)

**Selection Rate Ratio / 80% Rule**

Compares the positive prediction rate of a demographic group against the highest-selected group. Under the US EEOC 4/5ths (80%) rule, if the highest group is selected 80% of the time, all other groups should be selected at least 64% of the time (80% of 80%). A ratio below 0.80 indicates adverse impact.

`Ideal = 1.00; Safe threshold ≥ 0.80 • EEOC 4/5ths Rule / EU AI Act Annex IV §2(g)`

### Demographic Parity Difference (DPD)

**Positive Rate Percentage Gap**

Measures whether the model favors one group over another in raw volume of positive predictions. A difference of 0.15 means individuals from one group are 15 percentage points more likely to receive a positive prediction than individuals from another group.

`Ideal = 0.00; Safe threshold ≤ 0.10 • EU AI Act Art. 10(2)(f) / NIST AI RMF Measure 2.11`

### Equalized Odds Difference (EOD)

**Error Rate Equality Gap**

Ensures the model is equally accurate for all demographics. It verifies that qualified individuals have an equal chance of being accepted, and unqualified individuals have an equal chance of being rejected, regardless of identity.

`Ideal = 0.00; Safe threshold ≤ 0.10 • EU AI Act Art. 10(2)(g) / NIST AI RMF Measure 2.11`

### Equal Opportunity Difference (EOP)

**True Positive Opportunity Gap**

Focuses strictly on true positive outcomes (e.g. qualified applicants or correct identity matches). A difference of 0.08 means qualified individuals in one group are 8 percentage points less likely to be correctly recognized.

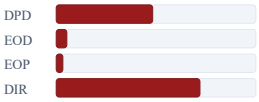
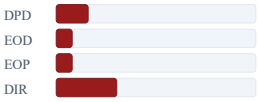
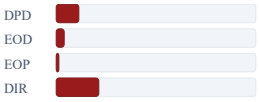
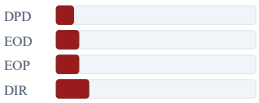
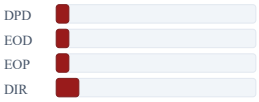
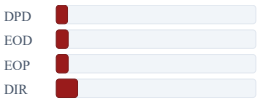
`Ideal = 0.00; Safe threshold ≤ 0.10 • EU AI Act Art. 10(3) / NIST AI RMF Measure 2.11`

| Metric Name | Point Estimate | 95% Bootstrap CI | CI Band vs. Null | Significance (p-value) | Regulatory Clause |
|---|---|---|---|---|---|
| **Demographic Parity Difference** | 0.192 | [0.156, 0.222] |  | p=0.0000 **p < 0.05** | EU AI Act Art. 10(2)(f) / NIST Measure 2.11 |
| **Equalized Odds Difference** | 0.053 | [0.035, 0.068] |  | p=0.0000 **p < 0.05** | EU AI Act Art. 10(2)(g) / NIST Measure 2.11 |
| **Equal Opportunity Difference** | 0.053 | [0.034, 0.069] |  | p=0.0000 **p < 0.05** | EU AI Act Art. 10(3) / NIST Measure 2.11 |

| Metric Name | Point Estimate | 95% Bootstrap CI | CI Band vs. Null | Significance (p-value) | Regulatory Clause |
|---|---|---|---|---|---|
| **Disparate Impact Ratio** | 0.713 | [0.676, 0.760] |   | p=0.0000 **p < 0.05** | EEOC 4/5ths Rule / EU AI Act Annex IV §2(g) |

## 2. Per-Subgroup Disparity Matrix

Detailed stratification across individual demographic groups, ordered by largest deviation from the null (no-disparity) point, most significant first. Subgroups below the NFR-003 sample threshold (n < 30) are strictly flagged, uncomputed to prevent small-sample estimation artifacts, and summarised separately below.

| Subgroup | Sample Size (n) | DPD | EOD (Max-Gap) | EOP | DIR (Symmetric) | Deviation from Null (red = p < 0.05) | Data Integrity Status |
|---|---|---|---|---|---|---|---|
| **Middle Eastern** | 1209 | 0.145 | 0.016 | 0.010 | 0.784 |   | **Audited (Valid)** |
| **Latino_Hispanic** | 1623 | 0.048 | 0.024 | 0.024 | 0.909 |   | **Audited (Valid)** |
| **Indian** | 1516 | 0.034 | 0.012 | 0.004 | 0.936 |   | **Audited (Valid)** |
| **Black** | 1556 | 0.026 | 0.034 | 0.034 | 0.951 |   | **Audited (Valid)** |
| **White** | 2085 | 0.019 | 0.019 | 0.019 | 0.966 |   | **Audited (Valid)** |
| **East Asian** | 1550 | 0.017 | 0.018 | 0.018 | 0.968 |   | **Audited (Valid)** |
| **Southeast Asian** | 1415 | 0.013 | 0.009 | 0.000 | 0.975 |   | **Audited (Valid)** |

## 3. Governance Documentation & Intended Use

> **EU AI Act (Regulation EU 2024/1689) Article 10 / Annex IV Alignment:** This audit dossier provides verifiable diagnostic documentation of demographic performance disparities. It does not perform model retraining or in-situ weight modifications.

**Model Description:** ImageNet-pretrained ResNet-34 terminating in an 18-unit linear multi-task head sliced for race (7), gender (2), and age (9) classification.

**Backend configuration:** dual. Both backend executions completed.

Cross-backend agreement checks point estimates only; bootstrap, hypothesis tests, and reporting share implementation. Consult execution logs for backend failures and divergences.